2×2 strategic calculation exam

7 systems • 5,184 cases • 29,952 individual answers per system

Complete-case accuracy

Every answer in a case must be correct

	Best responses	Strict dominance	Pure Nash equilibria	Level-k (k = 1–3)
Marres kernel + adapter	100%	100%	100%	100%
GPT-oss-120B · high reasoning	100%	100%	99.83%	94.21%
TypeSafe Jev	30.90%	37.50%	25.87%	18.46%
Torchcast Decision 27B	19.10%	41.84%	45.66%	14.70%
AutoTrust JEV-Gemma4-26B-A4B	7.47%	33.85%	37.33%	8.04%
Wald v2.1 · one-pass	0.35%	27.43%	16.67%	8.48%
Drex DLM	0%	13.19%	0.17%	1.85%

Individual-answer accuracy

Partial credit: each answer scored separately

	Best responses	Strict dominance	Pure Nash equilibria	Level-k (k = 1–3)
Marres kernel + adapter	100%	100%	100%	100%
GPT-oss-120B · high reasoning	100%	100%	99.96%	96.27%
TypeSafe Jev	86.39%	80.03%	70.70%	61.65%
Torchcast Decision 27B	81.53%	82.64%	83.72%	50.25%
AutoTrust JEV-Gemma4-26B-A4B	66.30%	79.25%	77.43%	42.76%
Wald v2.1 · one-pass	61.50%	76.48%	70.31%	45.92%
Drex DLM	49.80%	62.98%	28.52%	3.09%

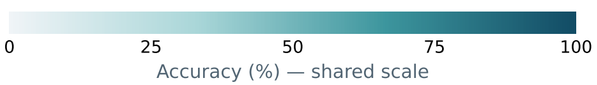

SCOPE: MATHEMATICAL CORRECTNESS ONLY

Tests explicitly specified best responses, strict dominance, pure Nash equilibria and level-k calculations.

Does not evaluate the full Marres application, real-world decision quality, or general intelligence.

Not an equal-compute comparison: Marres implements these calculations directly; GPT-oss used high reasoning (~9.4M reasoning tokens).

Invalid responses count as incorrect: GPT-oss 2; TypeSafe Jev 29. No missing cases. Frozen exam v2 • Runs: 16 October 2026.